Sanket Navnath Janger

Willing to Relocate | (609) 302-1913 | sjanger@binghamton.edu | LinkedIn | Github Link | Portfolio Link

EDUCATION

Binghamton University, State University of New York, Thomas J. Watson College of Engineering and Applied Science
Master of Science (M.S) in Computer Science (AI Track) — *Expected December 2026*
Savitribai Phule Pune University, Nashik, India
Bachelor of Engineering (B.E.) in Computer Engineering | Honors in Data Science — *August 2020 - June 2024*

TECHNICAL SKILLS

Languages: Python, JavaScript, TypeScript, SQL
Datastores: PostgreSQL, MySQL, Redis, BigQuery, ChromaDB
Frameworks & Libraries: FastAPI, Node.js, Express, React, Flask, SQLAlchemy, Alembic, Prisma
AI / LLM Systems: LLM APIs (Groq LLaMA 3.3), RAG, Embeddings, Vector Search (ChromaDB), FinBERT, sentence-transformers
Messaging/Streaming: Apache Kafka, RabbitMQ, Redis Pub/Sub, WebSockets
Cloud & DevOps: AWS (Cognito, S3, IAM), GCP (Pub/Sub, Cloud Run, Cloud Functions, BigQuery), Docker, Kubernetes, CI/CD, Git, Linux, Cloudflare (Workers, Durable Objects, D1)
Dev Tools: GitHub Co-Pilot, Cursor, Claude Code
Practices: Object-Oriented Programming (OOP), Regression Testing, Agile, Data Structures and Algorithms

PROFESSIONAL EXPERIENCE

Altheros Capital, LLC | *Los Angeles, CA, United States*
Back-End Lead Intern — *Sept 2026 - Present*

- Secured a fully unprotected API by building shared Cognito JWT middleware (RS256, JWKS validation) across all borrower, lender, financial, and user routes, so unauthenticated or tampered requests return 401 before reaching business logic.
- Led backend development for Creddly AI, a private-credit marketplace, directing 11 interns across authentication, role-based access control, KYC verification, AI matching, and payments.
- Reviewed intern pull requests before merge, enforcing ownership checks and input validation across borrower and lender endpoints.
- Designed a lender-borrower matching engine combining structured scoring, pgvector similarity, asynchronous SQS/Lambda recalculation, and Redis caching, then scoped it into GitHub issues with acceptance criteria.
- Planned a five-phase production roadmap and resolved an Alembic migration defect caused by a UTCDateTime/TIMESTAMP mismatch, eliminating false-positive schema changes.

Back-End Developer Intern — *June 2026 - August 2026*

- Remediated a broken object-level authorization vulnerability on a public endpoint and enforced field-level access controls across more than five profile endpoints.
- Implemented investor portal authorization with three permission tiers across four file categories, combining JWT claims and PostgreSQL queries to enforce document-level visibility.
- Developed Cara AI REST APIs using FastAPI, PostgreSQL, Redis, and Docker; wrote a 38-test regression suite that caught a foreign-key cascade bug before production.

PROJECT EXPERIENCE

OrderFlow: A Distributed Order Processing System | GitHub Link

- Achieved 99.7% fulfillment across 1,000 concurrent orders by claiming work with PostgreSQL row-level locking to prevent double processing, routing the remaining 0.3% to a RabbitMQ dead-letter queue.
- Diagnosed Kubernetes CPU-based autoscaling as the wrong signal for an I/O-bound workload after it cut workers from four to one mid-test, then added a minimum replica floor that recovered throughput from 0.36 to 1.37 orders per second (3.8x).
- Pushed order status changes to clients in real time through Redis Pub/Sub and WebSockets, so tracking needed no polling.

GrindMate: AI-Native Data Structures and Algorithms Tracker | Project Live Link

- Shipped natural language problem logging on Cloudflare's edge, measured by isolated per-user state and multi-tenant data safety, by building an AI-native DSA tracker with Llama 3.3, per-user Durable Objects, and D1 composite primary keys.
- Delivered spaced repetition reviews across the NeetCode 150 roadmap, measured by accurate 1/3/7-day review intervals via DO Alarms, by building a TypeScript React dashboard with Recharts visualizations across 18 categories and daily email reminders via Resend.
- Acquired 8+ organic users within 24 hours, measured by public sign-up traction, by launching GrindMate publicly at grindmate.dev.

FinPulse: Real-Time Financial News Intelligence Platform | GitHub Link

- Sustained 500+ articles per hour at sub-50ms API latency by decoupling ingestion from serving with Kafka (KRaft mode) and caching hot reads in Redis.
- Eliminated per-token sentiment API costs by routing classification to a local FinBERT model (22ms inference) and reserving Groq LLaMA 3.3 for summarization only.
- Enabled natural language queries by embedding articles with sentence-transformers and indexing them in ChromaDB, returning results in under 200ms across 1,000+ articles.