

# Sanket Navnath Janger

☎ +1 (609) 302-1913 | ✉ sjanger@binghamton.edu | 🔗 LinkedIn | 🐙 GitHub | Binghamton, NY

## Education

### State University of New York at Binghamton

Master of Science in Computer Science

*Binghamton, NY*

*Expected Graduation - December 2026*

### Savitribai Phule Pune University

B.E. in Computer Engineering (Honors: Data Science)

*Pune, India*

*Aug 2020 – Jul 2024*

## Technical Skills

**Languages:** Python, TypeScript, JavaScript, SQL, HTML/CSS

**Frameworks and Libraries:** FastAPI, Node.js, Express, React, Flask, SQLAlchemy, Alembic, Prisma

**Datastores:** PostgreSQL, MySQL, Redis, BigQuery, ChromaDB

**Messaging/Streaming:** Apache Kafka, RabbitMQ, Redis Pub/Sub, WebSockets, REST APIs

**Cloud & DevOps:** AWS (S3, IAM), GCP (Pub/Sub, Cloud Run, Cloud Function, BigQuery), Docker, Kubernetes, Cloudflare (Workers, Durable Objects, D1), Git, Linux

**AI Developer Tools/NLP:** RAG, LLM APIs, FinBERT, sentence-transformers, Codex, Cursor, Claude Code

**Networking:** TCP/IP, Routing Fundamentals

## Experience

### Back-End Developer Intern — Altheros Capital

*Los Angeles, CA June 2026 – August 2026*

- Delivered investor portal backend in 8 days against a 2-week estimate: migrated auth from AWS Cognito to JWT to remove external service dependency, built three-tier document visibility across 4 file categories (global, project-specific, investor-specific) and AWS S3 document management, using Node.js, Express, PostgreSQL, and Prisma
- Identified and patched a broken object-level authorization (BOLA) vulnerability on a public API endpoint, preventing cross-user data exposure and enforcing field-level access controls across 5+ authenticated profile endpoints using FastAPI and PostgreSQL
- Built RESTful APIs for Cara AI, a production AI companion platform, using FastAPI, PostgreSQL, Redis, and Docker, covering authentication, user profiles, settings, and session management with device and IP tracking.

## Academic Projects

### OrderFlow - Distributed Order Processing System

[GitHub Link](#)

- Built a distributed order processing system with RabbitMQ and PostgreSQL row-level locking, achieving 99.7% fulfillment across 1,000 concurrent orders, with the 0.3% correctly routed to the dead letter queue after exhausting payment retries
- Scaled workers from 1 to 4 replicas for 2.6x throughput improvement and added real-time order tracking via Redis pub/sub and WebSocket
- Deployed to Kubernetes with HPA, diagnosed CPU-based autoscaling as wrong for I/O-bound workloads after it silently scaled workers from 4 to 1 mid-load-test, fixed with a minReplicas floor and recovered 3.8x throughput versus the degraded single-worker state

### GCP Usage Analytics Pipeline

[GitHub Link](#)

- Built event-driven data pipeline ingesting UI interactions and API metrics through Pub/Sub into BigQuery via serverless Cloud Function, achieving under 2 seconds end-to-end latency from publish to query
- Reduced BigQuery scan volume via date partitioning and component-level clustering, so queries filter to relevant partitions instead of scanning the full table
- Deployed FastAPI REST layer on Cloud Run with error rate and latency alerting, backed by a live dashboard showing real-time event volume and per-endpoint latency. Integrated with GrindMate via a secured HTTP ingest endpoint, enabling real user events from a Cloudflare Workers app to flow into the same pipeline

### FinPulse - Real-Time Financial News Intelligence Platform

[GitHub Link](#)

- Architected real-time news intelligence platform processing 500+ articles/hour using Apache Kafka (KRaft mode), FastAPI, and WebSocket; achieved sub-50ms API latency with Redis caching and PostgreSQL persistence
- Implemented hybrid AI pipeline combining local FinBERT sentiment classification (20ms inference) with Groq LLaMA 3.3 for summarization; eliminated per-token API costs by running sentiment analysis locally
- Built RAG-based semantic search using sentence-transformers embeddings and ChromaDB, enabling natural language queries like "Fed impact on tech stocks" with under 200ms retrieval across 10K+ articles

### GrindMate

[Live Link](#)

- Built AI-native DSA tracker on Cloudflare's edge with natural language problem logging via Llama 3.3, per-user Durable Objects for isolated state, and D1 with composite primary keys for multi-tenant data safety
- Shipped spaced repetition reviews at 1/3/7-day intervals via DO Alarms, NeetCode 150 roadmap across 18 categories with FAANG tags, and daily email reminders via Resend. Acquired 8+ organic users within 24 hours